

COMP70068 Scheduling and Resource Allocation

Clemen Kok

December 2024

1 Introduction to Scheduling

What is Scheduling? Job - Unit of work we want to process; Machine - Resource that can process jobs. Sequence - When to start, end, and resume. Also consider parallel workflows, Due dates, Uncertain information.

Applications of Scheduling. Comp Arch, Cloud, Production Scheduling, Project Planning.

2 Single Machine Scheduling

Gantt Chart. How many schedules can we find to sequentially run n jobs without idling times on a single machine - $n!$

Sequencing. Optimising functions $f(C_1, \dots, C_n)$ - *regular measures* - monotonically non-decreasing with respect to each of the job completion times $C_j, j \leq n$.

Examples of regular measures:

1. $C_{max} = \max(C_1, \dots, C_n)$ - time at which the last operation is completed (Makespan)
2. $\sum C_j$ - Total completion time
3. $\sum w_j C_j$ - total weighted completion time
4. $\sum w_j (1 - e^{-rC_j})$ - discounted weighted completion time with discount rate $0 \leq r \leq 1$ (e.g. inflation)

Why would you choose to minimise C_{max} ? Minimises Makespan, Total time needed to complete all jobs, Ensures set of jobs is completed ASAP, Ensures system resources are utilised efficiently, maximises throughput, higher cost efficiency, Minimise idleness.

When to NOT minimise C_{max} ? - Cost to idle.

Classifying Scheduling Problems. Use notation $\alpha|\beta|\gamma$ with α : machines, β : job characteristics and γ : optimality criterion.

Assignments:

1. $\alpha = 1$ Single Machine, P Parallel Machines
2. β : *pmtn* - preemption (interrupts and suspension for higher priorities), r_j : release times (time when job becomes available for scheduling), d_j : due dates, *prec* precedences
3. γ : C_{max} , ΣC_j , $\Sigma w_j C_j$
4. E.g. $1||\Sigma C_j$, $1|r_j, pmtn|\Sigma C_j$, $P|prec|\Sigma w_j C_j$

PARTITION. Theoretical problem that illustrates the hardness of scheduling for load balancing with known, deterministic, job processing times. Given a set of S integers (e.g. job processing times) with repetitions allowed, **PARTITION** asks to form two subsets with identical sum, or return that they do not exist. *NP*-hard. With $S = \{1, 1, 1, 2, 2, 2, 3\}$ we may choose $S_1 = \{1, 1, 1, 3\}$ and $S_2 = \{1, 2, 3\}$ or ... or ...

Proof that PARTITION is NP Complete. Reduce from Subset Sum Problem.

3 Release Times

Assumptions

- A set of n jobs is immediately ready to be scheduled.
- Deterministic Processing Times, known in advance.
- No preemption.

- Single server machine.
- No explicit setup times.
- No machine idling while jobs are waiting.
 - Idle - non pre-emptive - higher priority.

Preemptive vs Non-preemptive Scheduling

- No interrupts can lead to improved C .
- In single machine scheduling, if performance is monotonic with respect to completion times, there is no advantage in using preemption.

Scheduling on a Single Resource

- Scheduling on a single resource is the simplest scheduling problem.
- More complex real-world situations can often reduce to analyzing single machine problems when we model the *bottleneck* machine.
- Non-idling schedules: no gains in regular measures by inserting idle times within a single machine. Not true in general outside assumptions.
- $1||C_{\max}$ is trivial: under non-idling, makespan minimization is optimized by any random schedule since:

$$C_{\max} = \sum_j p_j = \text{const}$$

where p_j is the processing time of job $j = 1, \dots, n$.

Total Completion Time Problem ($1||\sum C_j$)

- Shortest Processing Time (SPT) Rule is optimal.
- Let $[j]$ denote the j -th scheduled job, which will have processing time $p_{[j]}$ and completion time $C_{[j]}$. Since all jobs are released at time 0, the latter will be:

$$\begin{aligned} C_{[1]} &= p_{[1]} \\ C_{[2]} &= p_{[1]} + p_{[2]} \\ &\vdots \\ C_{[n]} &= p_{[1]} + p_{[2]} + \dots + p_{[n]} \end{aligned}$$

$$\sum C_j = np_{[1]} + (n-1)p_{[2]} + \dots + p_{[n]}.$$

- Shorter jobs are best scheduled earlier.
- Note: job completion time is accumulative.

Q: Computational Complexity of SPT? $P|n \log n$

1|| $\sum w_j C_j$ **Problem**

- Weighted Shortest Processing Time (WSPT) Rule (Smith's rule).
- Schedules jobs in non-decreasing p_j/w_j ratio order.
- Proof follows adjacent pairwise interchange argument:
 - Assume that S is optimal, but not a WSPT schedule. Thus, we can find two adjacent jobs i and j in S such that $\frac{p_i}{w_i} > \frac{p_j}{w_j}$.

Proof of WSPT Optimality

$$\begin{aligned} \sum_{l=1}^n w_l C_l &= w_i (p(B) + p_i) + w_j (p(B) + p_i + p_j) + \sum_{k \neq i, j} w_k C_k \\ \sum_{l=1}^n w_l C'_l &= w_j (p(B) + p_j) + w_i (p(B) + p_i + p_j) + \sum_{k \neq i, j} w_k C_k \\ \sum w_l C_l - \sum w_l C'_l &= w_j p_i - w_i p_j > 0 \end{aligned}$$

Since $S - S' > 0$, $S > S'$.

Key Idea: Swapping two jobs in the sequence affects completion time proportional to weights and processing times.

Job Chains as an Extension of WSPT

- Sequential precedences among jobs.
- Non-interruptible execution.
- Let:

$$\begin{aligned} A : 1 &\rightarrow 2 \rightarrow \dots \rightarrow k \\ B : k+1 &\rightarrow k+2 \rightarrow \dots \rightarrow n. \end{aligned}$$

Which should be processed first?

- A should be scheduled before B if:

$$\frac{\sum_{j=1}^k p_j}{\sum_{j=1}^k w_j} < \frac{\sum_{j=k+1}^n p_j}{\sum_{j=k+1}^n w_j}$$

- Similar to WSPT (non-decreasing chain ratios).

Interruptible Execution:

- Interleave executions of jobs from different chains.
- Individual jobs are still executed non-preemptively.

- WSPT may still be applied to sub-chains.

ρ -factor:

$$\rho_A = \min_{1 \leq l \leq k} \frac{\sum_{j=1}^l p_j}{\sum_{j=1}^l w_j}.$$

- Initial subchains of A of length l : the one with the minimum ratio determines the ρ factor.

Uncertain Processing Times

- Consider $1||\sum C_j$ and $1||\sum w_j C_j$ with uncertain p_i .
- Relaxes assumption A2.
- Assume each processing time p_i is sampled from a statistical distribution, making p_i a random variable.
- Model assumption A5 (Machine Breakdowns): A breakdown could unexpectedly extend the processing time of the job in service until it resumes after machine repair.

Key Points:

- Maximum is not tractable in stochastic vs deterministic cases.
- Optimize expected total completion time $E[\sum C_j]$.
- SPT can be replaced by S-Expected-PT policy, scheduling jobs in order of mean processing time.

Let p_j be independent random variables:

$$E[\sum C_j] = E[n p_{[1]} + (n-1)p_{[2]} + \dots + p_{[n]}] = nE[p_{[1]}] + (n-1)E[p_{[2]}] + \dots + E[p_{[n]}].$$

If p_j are exponentially distributed with rates $\mu_j = 1/E[p_j]$ and weights are c_j , the principle is known as the $c\mu$ rule.

4 Due Dates

Key Definitions

Each job has a due date d_j (delivery time). Due dates are soft and can be violated. The goal is to minimise deviations. The following metrics are defined:

- **Lateness:** $L_j = C_j - d_j$
- **Tardiness:** $T_j = \max(0, C_j - d_j)$

Minimising Total Lateness ($1||\Sigma L_j$)

Although the Shortest Processing Time (SPT) rule ignores the due dates d_j , it remains optimal for minimising total lateness.

$$\Sigma_j L_j = \Sigma_j (C_j - d_j) = \Sigma_j C_j - \Sigma_j d_j$$

Since $D = \Sigma_j d_j$ is given and independent of the schedule, minimising total lateness is equivalent to minimising total completion time. A similar situation arises for the weighted analogue $1||\Sigma w_j L_j$.

Focusing on L_{\max}

The total lateness ΣL_j does not fully capture the significance of due dates. Hence, the problem $1||\Sigma L_j$ is seldom studied. Instead, the focus is often on minimising the maximum lateness L_{\max} :

$$L_{\max} = \max_j L_j$$

Earliest Due Date (EDD) Rule

For jobs with different due dates $d_i \geq d_j$, the **Earliest Due Date (EDD)** rule, which schedules jobs in non-decreasing d_j order, is optimal for $1||L_{\max}$. Ties are broken at random.

EDD is also optimal for $1||T_{\max}$. The relationship between T_{\max} and L_{\max} is:

- $T_{\max} = 0$ when $L_{\max} \leq 0$
- $T_{\max} = L_{\max}$ if $L_{\max} > 0$

Proof of EDD Optimality

Proof details can be found in the notes.

Stochastic Case

Under uncertainty, lateness and tardiness are harder to handle due to expectations involving maxima of random variables. However, it can still be shown that EDD remains optimal for $1||L_{\max}$ and $1||T_{\max}$.

Minimising Total Tardiness ($1||\Sigma T_j$)

Minimising ΣT_j provides more control over tardiness distribution across jobs. However:

- The problem is NP-hard.

- No optimal rule like EDD or SPT exists.
- EDD is still optimal in restricted cases.

Agreeable Jobs

Jobs are **agreeable** if, for any two jobs i and j , $p_i \geq p_j$ implies $d_i \geq d_j$. If all jobs are agreeable, total tardiness is minimised by EDD sequencing with ties broken by SPT. This can be proven using adjacent pairwise interchange analysis.

Minimising Number of Tardy Jobs ($\sum_j U_j$)

The number of tardy jobs is defined as:

$$U_j = \begin{cases} 0 & \text{if } C_j \leq d_j \\ 1 & \text{otherwise} \end{cases}$$

Moore-Hodgson Algorithm

The problem $1||\sum_j U_j$ is solvable exactly in polynomial time using the **Moore-Hodgson Algorithm**. However, the weighted case is NP-hard.

1. Jobs $j = 1, \dots, n$ are added one by one, in EDD order, to the on-time schedule S_O .
2. If job j completes after time d_j , the job $i \in S_O$ with the largest p_i is declared late and moved to the late schedule S_L in an arbitrary position.
3. Step 2 is repeated until either job j can be added to S_O or it is moved to S_L if it has the largest p_i .
4. After all jobs are added, the algorithm returns $S = \{S_O, S_L\}$:
 - On-time jobs (S_O) are scheduled before late jobs (S_L).
 - The order of jobs in S_L is arbitrary.

Example:

- Remove J_2 , then J_1 (by processing time p_i).
- These jobs will be late regardless.

5 Enumeration and Local Search

Introduction

Many scheduling problems are NP-Hard and cannot be solved efficiently if the problem size is large enough. Combinatorial optimization involves minimizing a cost function $g(S)$ over all feasible schedules S .

In most cases, the cost function is an additive function:

$$g(S) = \sum_j g_j(S),$$

where $g_j(S)$ is the cost of scheduling job j prescribed by schedule S . - Example: $g_j(S) = w_j U_j$ defines a $1||\sum_j w_j U_j$ problem. - Example: $g_j(S) = \max(0, C_j - d_j)$ defines a $1||\sum_j T_j$ problem.

Local Search

Similar to hill climbing but applied to discrete problems. A neighborhood $\mathcal{N}(S)$ of the current solution S is generated using adjacent pairwise interchange, subject to constraints. The solution evolves by selecting the next one in $\mathcal{N}(S)$. May get trapped in local optima.

Simulated Annealing (SA)

A random-walk-based search that explores the solution space while being guided towards the global optimum over time. Escape tendency around local optima is controlled by a temperature parameter T_k , reduced over time until convergence.

Algorithm:

1. Start with an initial solution x_0 .
2. At iteration $k + 1$, randomly choose a neighbour $y \in \mathcal{N}(x_k)$.
3. Compute cost change $\Delta = g(x_k) - g(y)$: If $\Delta \geq 0$, accept y as $x_{k+1} = y$. If $\Delta < 0$, accept y with probability e^{Δ/T_k} .
4. Update temperature:

$$T_{k+1} = \alpha T_k = \alpha^{k+1} T_0, \quad T_0 > 0, \quad 0 < \alpha < 1.$$

5. Stop when $k > K$.

Alternative cooling model:

$$T_k = \frac{T_0}{1 + \log(1 + k)}.$$

Tabu Search

A deterministic global optimization method using a memory structure called the tabu list \mathcal{T} . Tracks recently swapped job pairs to avoid cycling back to previous solutions. Exceptions (aspiration criteria) allow revisiting a solution if it improves the best overall result.

Algorithm:

1. Swap adjacent jobs i and j in schedule S . Add pair (i, j) to \mathcal{T} .
2. Prevent re-swapping (i, j) for L iterations unless aspiration criteria are met.
3. Stop when no further improvements are possible.

Handling Constraints

Schedules may have equality and inequality constraints:

$$g_{\text{best}} = \min_S g(S), \quad \text{s.t. } f_i(S) = 0, \quad h_k(S) \leq 0.$$

Rejection Method

Skip candidate solutions violating constraints. Valid initial schedules required.

Penalty Function Method

Modify the cost function:

$$g'(S) = g(S) + \sum_{i=1}^I \lambda_i f_i(S) + \sum_{k=1}^K \gamma_k \max(0, h_k(S)).$$

Dynamic Programming (DP)

Define J as a set of jobs. Let $G(J)$ be the minimum cost of scheduling the jobs in J , where $G(J)$ is obtained recursively:

$$G(J) = \min_{j \in J} \{G(J - \{j\}) + g_j(J)\}, \quad G(0) = 0.$$

Complexity: $O(n2^n)$.

Branch and Bound

Combines branching (partitioning solutions) and bounding (eliminating suboptimal solutions). Nodes in a search tree represent subproblems with specific jobs fixed in the schedule. Subproblems are pruned using bounds to save computation.

6 Parallel Machine Scheduling

Introduction

1. Jobs are mapped to m identical parallel machines. 2. A job cannot run on two machines at the same time. 3. The makespan problem, denoted $P||C_{max}$, is no longer trivial and equates to balancing job loads across machines. 4. For $m = 2$, the problem is denoted P_2 (parallel scheduling on two machines).

Total Completion Time Problems

1. The Shortest Processing Time (SPT) optimality for $1||\sum C_j$ can be generalized to $P||\sum C_j$. 2. Adding indices to represent completion times and processing on m machines, the goal is to minimize:

$$\begin{aligned} \sum_{j=1}^{n_1} C_{1,j} &= n_1 p_{1,[1]} + (n_1 - 1)p_{1,[2]} + \cdots + p_{1,[n_1]}, \\ &\vdots \\ \sum_{j=1}^{n_m} C_{m,j} &= n_m p_{m,[1]} + (n_m - 1)p_{m,[2]} + \cdots + p_{m,[n_m]}. \end{aligned}$$

3. SPT optimizes the system by sorting jobs in ascending size and then applying round-robin across machines. 4. Weighted Shortest Processing Time (WSPT) assigns jobs to the first idle machine in increasing p_i/w_i ratios. While not optimal for $P||\sum w_j C_j$, which is NP-hard, it serves as an approximation with a worst-case ratio:

$$R = \frac{1 + \sqrt{2}}{2} \approx 1.21.$$

5. Round-robin schedules jobs cyclically through m_1, m_2, \dots, m_n .

Makespan Problems

1. With m machines, the makespan admits the lower bound:

$$M^* = \max(p_{max}, \frac{\sum p_j}{m}),$$

where $p_{max} = \max p_j$. 2. The term $\sum p_j/m$ represents the optimal case where all machines finish jobs simultaneously. 3. The p_{max} term arises because no job can run on multiple machines simultaneously, so the longest job determines a minimum makespan.

McNaughton's Wrap Around Rule

1. Schedule an arbitrary job on machine m_1 at time 0. 2. Start any unscheduled job as soon as possible on the same machine. 3. Repeat the process until the makespan on a machine exceeds M^* or all jobs are scheduled. 4. Reassign processing beyond M^* to the next machine, starting at time 0. Repeat the process.

Non-Preemptive Case

1. The non-preemptive makespan problem is NP-hard, even with $m = 2$ machines. 2. Assuming integer p_j and $\sum p_j$ divisible by 2, $P_2||C_{max}$ asks to partition integers into two subsets summing to M^* , equivalent to the partition problem.

List Scheduling (LS)

1. List scheduling is a popular approximation. Given a priority list of jobs, LS assigns the first available job to the first available machine. 2. The worst-case ratio for $P||C_{max}$ is:

$$R = 2 - \frac{1}{m}.$$

Proof of the Worst-Case Ratio

1. Let $t_k = C_{max}^{LS} - p_k$ be the time when job k starts. Since machines are always busy before t_k :

$$m \cdot t_k \leq \sum_{j=1}^n p_j - p_k.$$

2. Substituting $t_k = C_{max}^{LS} - p_k$ and rearranging:

$$C_{max}^{LS} \leq \frac{\sum_{j=1}^n p_j}{m} + \frac{(m-1)p_k}{m}.$$

3. The optimal makespan satisfies $C_{max}^{OPT} \leq M^*$, where:

$$\frac{\sum_{j=1}^n p_j}{m} \leq C_{max}^{OPT}, \quad p_k \leq p_{max} < C_{max}^{OPT}.$$

4. Therefore:

$$\frac{C_{max}^{LS}}{C_{max}^{OPT}} = 2 - \frac{1}{m}.$$

Longest Processing Time (LPT)

1. List scheduling achieves a $(2 - \frac{1}{m})$ approximation if no specific job priority order is specified. 2. LPT assigns jobs in non-increasing order of size. LS with LPT achieves a smaller worst-case ratio:

$$R = \frac{4}{3} - \frac{1}{3m}.$$

7 Workflow Scheduling

Directed Acyclic Graphs (DAGs)

1. Job precedences are represented as a Directed Acyclic Graph (DAG) $G = (V, E)$:
 - Vertices V are jobs.
 - An edge $(i, j) \in E$ indicates that job i must complete before job j can start.
2. Precedence constraints frequently require forced idleness.
3. Scheduling theory under precedence constraints is challenging:
 - Most problems with precedences are NP-hard.
 - Complexity remains an open question in some cases.
4. Optimal methods exist only under restrictive assumptions:
 - Unit processing times ($p_j = 1, \forall j$).
 - Fixed number of processors (e.g., two).
 - Restricted topologies (chains, in-trees, out-trees):
 - In-tree: Every node has at most one child.
 - Out-tree: Every node has at most one parent.
 - Chain: Every node has at most one parent and one child.

Hu's Algorithm

1. Hu's algorithm solves $P|n, p_j = 1|C_{max}$ (parallel scheduling with in-tree precedence constraints).
2. Assign a level α_j to each job j as follows:
 - The exit node is labeled with $\alpha_j = 1$.
 - For all other nodes:

$$\alpha_j = 1 + \max_{i: (j,i) \in E} \alpha_i.$$

3. The maximum level is $L = \max_i \alpha_i$, and $\alpha_j - 1$ is the path length from j to the exit node. The longest path is called the critical path.
4. Hu's algorithm schedules ready jobs in non-decreasing α_j order, with complexity $\mathcal{O}(n)$.
5. The algorithm achieves optimal schedules for $P|i.t., p_j = 1|C_{max}$.

Optimality of Hu's Algorithm

1. Reduction to in-tree cases ensures optimality:
 - Reverse the arc orientations in out-tree problems to create an equivalent in-tree problem.
 - For disjoint in-trees, add a dummy task as the successor of all exit nodes.
2. The algorithm approximates $P|prec, p_j = 1|C_{max}$:
 - Worst-case ratio:

$$R = \begin{cases} \frac{4}{3}, & \text{if } m = 2, \\ 2 - \frac{1}{m-1}, & \text{if } m \geq 3. \end{cases}$$

Muntz-Coffman Algorithm

1. The Muntz-Coffman algorithm extends Hu's algorithm to the preemptive case.
2. A subset sequence S_1, S_2, \dots, S_k is defined as follows:
 - Each job $a \in G$ belongs to some S_i .
 - If $a \in S_j$ is a successor of $b \in S_i$, then $j > i$.
3. Given a subset sequence, schedule subsets in increasing i :
 - If $|S_j| > m$, use McNaughton's wrap-around rule, assigning $|S_j|/m$ time on m processors.
 - If $|S_j| \leq m$, assign one unit of time on $|S_j|$ processors.
4. The method is optimal for general DAGs, subject to $m = 2$ and $p_j = 1$.

Optimality of Muntz-Coffman Algorithm

1. The algorithm is optimal for:
 - $P_2|pmtn, prec, p_j = 1|C_{max}$.
 - $P|pmtn, i.t., p_j = 1|C_{max}$.

- $P|pmtn, o.t., p_j = 1|C_{max}$.
2. Under preemptive scheduling, $p_j > 1$ can be divided into jobs with $p_j = 1$.
 3. The worst-case ratio for general preemptive scheduling $P|pmtn, prec|C_{max}$ is:

$$R = 2 - \frac{2}{m}, \quad m \geq 2.$$
 4. The algorithm is optimal for $m = 2$ and near-optimal for small $m > 2$.

8 Bottleneck Analysis

System Overview

1. Jobs arrive at the system, are processed by machines, and then completed.
2. Assumptions:
 - A1: The system has M machines with arbitrary speeds.
 - A2: Jobs arrive at arbitrary times, not known in advance.
 - A3: Jobs may visit more than one machine and do not incur communication overheads.
 - A4: Job processing times are arbitrary and machine-dependent.
 - A5: Job sequencing at machines is arbitrary.
 - A6: Inside the system, there is no job creation, destruction, or parallel execution of a single job.
3. Key questions:
 - Bottleneck Analysis: Which machines limit the peak completion rate?
 - What-if Analysis: How will changes in job arrival rates or machine speeds impact the system?
4. Performance factors:
 - Arrival rate of jobs.
 - Processing time of jobs at the machine.
 - Utilization (fraction of time the machine is busy).
 - Contention.

Job Classes

1. The open system:

- Processes jobs using M machines.
- Offers C types of services (job classes).
- Receives class- c jobs at rate λ_c .

2. For a class- c job:

- Visits machine i , on average, v_{ic} times.
- Requires a mean processing time p_{ic} for each visit.
- Demand:

$$D_{ic} = v_{ic}p_{ic},$$

the total processing time accumulated on average during visits to machine i .

Operational Analysis

1. Monitor a machine for an observation period T seconds and collect:

- A_c : Total number of arrived jobs of class c .
- B_{ic} : Total time machine i is busy processing class- c jobs.
- F_c : Total number of finished jobs of class c .

2. Compute:

- $\lambda_c = A_c/T$: Average arrival rate of jobs of class c .
- $X_c = F_c/T$: Average system throughput of jobs of class c .
- $U_{ic} = B_{ic}/T$: Utilization of machine i for class- c jobs.

3. Stability:

- If the number of pending jobs remains finite, the system is stable.
- In a stable system:

$$\lambda_c = X_c = \lim_{T \rightarrow \infty} \frac{F_c}{T}.$$

- Unstable systems cannot cope with arrival rates, causing the backlog of pending jobs to grow unbounded.

Utilization Law

1. Demand:

$$D_{ic} = \frac{B_{ic}}{A_c}.$$

2. Utilization:

$$U_{ic} = \frac{B_{ic}}{T} = \frac{A_c}{T} \cdot \frac{B_{ic}}{A_c} = \lambda_c D_{ic}.$$

3. For C job classes, the total utilization of machine i is:

$$U_i = \sum_{c=1}^C U_{ic} = \sum_{c=1}^C \lambda_c D_{ic} = \sum_{c=1}^C X_c D_{ic}.$$

4. Estimating demand:

- Use multivariate linear regression to fit hyperplanes to samples of U_i and λ_c .
- Estimated demands D_{ic} are hardware-dependent and change after machine upgrades.

Bottleneck Analysis

1. Bottlenecks:

- Machines that limit performance by struggling to handle arrival rates.
- Tend to operate near 100

2. Single class ($C = 1$):

- For every machine i :

$$U_i = \lambda D_i.$$

- Maximum arrival rate:

$$\lambda \leq \frac{1}{D_{\max}},$$

where $D_{\max} = \max(D_1, \dots, D_M)$.

- Bottlenecks correspond to machines with D_{\max} .

3. Multi-class ($C > 1$):

- Machine usage is described by:

$$U_i = \sum_{c=1}^C \lambda_c D_{ic}.$$

- A machine j can saturate if there exists a mix of arrival rates $\lambda_1, \dots, \lambda_C$ such that $U_j = 1$.

- Verify saturation using linear programming (LP):

$$U_j^{\max} = \max \sum_{c=1}^C \lambda_c D_{jc},$$

subject to:

$$\sum_{c=1}^C \lambda_c D_{ic} \leq 1, \quad \lambda_c \geq 0.$$

- If $U_j^{\max} = 1$, the machine can saturate.

9 Competitive Decision Making

Introduction

1. Multi-agent settings lack a clear notion of what is optimal.
 - Focus is on equilibria rather than global optimizers.
 - An equilibrium is "good" if it overlaps with system optima.
2. Game setup:
 - Players: $i = 1, \dots, N$.
 - Each player has a set of actions $x_i \in \mathcal{X}_i$.
 - Each player incurs a cost $J_i(x_1, \dots, x_N)$, which depends on everyone's choices.
3. Example: The 2/3 game.

$$J_i(x_1, \dots, x_N) = \left| \frac{2}{3} \cdot \frac{1}{N} \sum_j x_j - x_i \right|.$$

4. Bimatrix representation example:

	R	S
R	(30, 0)	(30, 10)
S	(100, 100)	(0, 10)

Nash Equilibria

1. A Nash equilibrium (N.E.) models emergent behavior.
 - A pure N.E. is a feasible allocation such that no player can decrease their cost by unilateral deviation.

2. Definition:

$(x_1^*, \dots, x_N^*) \in \mathcal{X}_1 \times \dots \times \mathcal{X}_N$ is a N.E. if $J_i(x_i^*, x_{-i}^*) \leq J_i(x_i, x_{-i}^*) \quad \forall x_i \in \mathcal{X}_i, \forall i$.

Here:

- x_{-i} represents the strategies of all players except i .
- x_i is the strategy of player i .

3. Example:

	A	B
A	(30, 0)	(30, 10)
B	(11, 11)	(0, 10)

The N.E. is (10, 10), as neither player has an incentive to deviate.

Best Response Algorithm

1. Given other players' strategies x_{-i} , the best response is:

$$BR(x_{-i}) = \arg \min_{x_i \in \mathcal{X}_i} J_i(x_i, x_{-i}).$$

2. Steps:

- Initialize strategies x and set $i = 1$.
- Check if x is an equilibrium; if so, stop.
- Update $x_i \leftarrow BR(x_{-i})$.
- Set $i \leftarrow (i \bmod N) + 1$.
- Repeat.

3. Convergence:

- If the current allocation minimizes the cost, no one will change their strategy.
- The algorithm converges to a Nash equilibrium.

Mixed Nash Equilibria

- Players may randomize their actions, minimizing expected cost.
- For finite actions, let $\sigma_i \in \Delta_i$ represent the probabilities of player i 's actions. The expected cost is:

$$C_i(\sigma_1, \dots, \sigma_N) = \mathbb{E}_{x \sim \sigma} [J_i(x)] = \sum_{x \in \mathcal{X}} p(x) J_i(x).$$

3. Definition:

$(\sigma_1^*, \dots, \sigma_N^*) \in \Delta_1 \times \dots \times \Delta_N$ is a mixed N.E. if $C_i(\sigma_i^*, \sigma_{-i}^*) \leq C_i(x'_i, \sigma_{-i}^*) \quad \forall x'_i \in \mathcal{X}_i, \forall i$.

Matching Pennies Example

1. Game:

	H	T
H	$(1, -1)$	$(-1, 1)$
T	$(-1, 1)$	$(1, -1)$

2. No pure Nash equilibria exist.

3. Mixed Nash equilibria:

- Let p be the probability the row player chooses H .
- Let q be the probability the column player chooses H .

4. Expected payoffs:

$$\text{Row Player (P1): } \begin{cases} H : 1 \cdot q + (-1) \cdot (1 - q) = 2q - 1, \\ T : (-1) \cdot q + 1 \cdot (1 - q) = -2q + 1. \end{cases}$$

$$\text{Column Player (P2): } \begin{cases} H : -1 \cdot p + 1 \cdot (1 - p) = -2p + 1, \\ T : 1 \cdot p + (-1) \cdot (1 - p) = 2p - 1. \end{cases}$$

5. Solve for indifference:

$$2q - 1 = -2q + 1 \implies q = \frac{1}{2},$$

$$-2p + 1 = 2p - 1 \implies p = \frac{1}{2}.$$

6. The mixed Nash equilibrium is $(p, q) = (\frac{1}{2}, \frac{1}{2})$.

7. Expected cost:

$$\mathbb{E}[C_1] = \frac{1}{2} \cdot \frac{1}{2} - \frac{1}{2} \cdot \frac{1}{2} - \frac{1}{2} \cdot \frac{1}{2} + \frac{1}{2} \cdot \frac{1}{2} = 0.$$

Similarly, $\mathbb{E}[C_2] = 0$.

8. Deviation:

- If one player deviates, the expected cost remains 0.

10 Potential and Congestion Games

Potential Games

1. A potential game is a "nice" game where the cost functions of all players can be described by a single potential function.

2. Formal definition:

- A strategic game is a potential game if there exists a function $\Phi : \mathcal{X} \rightarrow \mathbb{R}$ such that:

$$J_i(x_i, x_{-i}) - J_i(y_i, x_{-i}) = \Phi(x_i, x_{-i}) - \Phi(y_i, x_{-i}),$$

for all x_i, y_i, x_{-i} , and for all players i .

3. The differences in the potential function, rather than its absolute values, determine the dynamics.

4. A potential game guarantees:

- Existence of at least one Nash Equilibrium (NE).
- Convergence of the Best Response (BR) algorithm in a finite number of steps for games with finite actions.

5. Intuition:

- Each time a player improves their utility, the potential changes accordingly.
- Cyclic behavior is impossible as it would imply $\Phi(x^1) > \Phi(x^2) > \dots > \Phi(x^1)$, which is contradictory.

Congestion Games

1. Congestion games model scenarios where players share resources, and the cost depends on resource usage.

2. Components:

- Set of resources \mathcal{R} .
- Resource cost functions $\uparrow_r(\cdot)$: latency functions specific to each resource.
- Set of players $\{1, \dots, N\}$.
- Feasible set $\mathcal{X}_i \subseteq 2^{\mathcal{R}}$ for each player i , representing the resources they can choose.
- Player cost:

$$J_i(x) = \sum_{r \in x_i} \uparrow_r(|x|_r),$$

where $|x|_r$ is the number of players using resource r .

3. Example: Load balancing.

- Resources are machines.
- Players are tasks managed by players.

- Each player's cost equals the runtime of their chosen machine.
4. Best Response (BR) convergence:
- Congestion games are potential games, with potential function:

$$\Phi(x) = \sum_{r \in \mathcal{R}} \sum_{j=1}^{|x|_r} \uparrow_r(j).$$

- The BR algorithm converges because the potential decreases with each step.

Convergence in Singleton Congestion Games

1. Definition:
 - A singleton congestion game is a congestion game where each player chooses a single resource ($|x_i| = 1$ for all $x_i \in \mathcal{X}_i$).
2. Example:
 - Load balancing games are singleton congestion games because each task can choose only one machine.
3. Convergence:
 - In singleton congestion games with n players and m resources, BR converges in $O(n^2m)$.
4. Proof intuition:
 - Replace original latencies with bounded integer values while preserving player preferences.
 - The potential $\Phi(x)$ measures total latency across all resources:

$$\Phi(x) = \sum_{r=1}^m \sum_{j=1}^{|x|_r} \bar{\uparrow}_r(j),$$

where $\bar{\uparrow}_r(j)$ are the sorted integer latencies for resource r .

- The potential is bounded by n^2m , as the maximum decrease in potential per BR step is at least 1.
5. Example calculation:
- Suppose R_1 and R_2 are resources with latencies:

$$\uparrow_{R_1}(1) = 1, \quad \uparrow_{R_1}(2) = 3, \quad \uparrow_{R_1}(3) = 5,$$

$$\uparrow_{R_2}(1) = 2, \quad \uparrow_{R_2}(2) = 4, \quad \uparrow_{R_2}(3) = 6.$$

After sorting latencies, assign new integer values:

$$\bar{\uparrow}_{R_1}(1) = 1, \quad \bar{\uparrow}_{R_2}(1) = 2, \quad \bar{\uparrow}_{R_1}(2) = 3.$$

- Compute potential:

- R_1 has 2 players. Sum latencies:

$$\bar{\downarrow}_{R_1}(1) + \bar{\downarrow}_{R_1}(2) = 1 + 3 = 4.$$

- R_2 has 1 player. Sum latency:

$$\bar{\downarrow}_{R_2}(1) = 2.$$

- Total potential:

$$\Phi(x) = 4 + 2 = 6.$$

- The potential decreases with each BR step and is bounded by n^2m .

11 Efficiency of Equilibria

Braess Paradox

1. Consider two routes from A to B :
 - North: road and ferry.
 - South: ferry and road.
2. Initially, 200 players split equally between the two routes (100, 100).
3. Introducing a new bridge changes the equilibrium:
 - Everyone avoids the ferry, resulting in:

$$J_i(x) = 2(15 + 0.1 \cdot 200) = 70, \quad \forall i.$$

- Using the road and ferry leads to a higher cost:

$$15 + 0.1 \cdot 200 + 400 = 75.$$

4. The new equilibrium increases overall travel cost, illustrating Braess Paradox.

Price of Anarchy (PoA)

1. Definition:
 - In a strategic game, a social cost function $SC : \mathcal{X} \rightarrow \mathbb{R}$ measures the quality of each allocation for the whole population.
 - The price of anarchy is defined as:

$$PoA = \frac{\max_{x \in PNE} SC(x)}{\min_{x \in \mathcal{X}} SC(x)}.$$

2. Interpretation:

- Measures the performance degradation from selfish decision-making.
- $SC(x)$ is typically the sum of players' costs.

3. Example:

- Worst Nash equilibrium (no ferry): $SC = 14000$.
- Optimal allocation: $SC = 12875$.
- Price of anarchy:

$$PoA = \frac{14000}{12875} \approx 1.087.$$

4. Implication:

- PoA quantifies the inefficiency caused by selfish behavior.
- Use incentives to reduce PoA .

Bounding the Price of Anarchy

1. Definition of a smooth game:

- A strategic game is (λ, μ) -smooth if there exist constants $\lambda > 0$ and $\mu < 1$ such that:

$$\sum_i J_i(x'_i, x_{-i}) \leq \lambda SC(x') + \mu SC(x), \quad \forall x', x \in \mathcal{X}.$$

2. PoA bound:

- In a (λ, μ) -smooth game:

$$PoA \leq \frac{\lambda}{1 - \mu}.$$

3. Proof:

- Let x be a Nash equilibrium and x' an optimal allocation.
- From the definition of Nash equilibrium:

$$SC(x) = \sum_i J_i(x) \leq \sum_i J_i(x'_i, x_{-i}).$$

- Using the smoothness property:

$$\sum_i J_i(x'_i, x_{-i}) \leq \lambda SC(x') + \mu SC(x).$$

- Combining:

$$SC(x) \leq \lambda SC(x') + \mu SC(x).$$

- Rearranging:

$$\begin{aligned} SC(x)(1 - \mu) &\leq \lambda SC(x'), \\ \frac{SC(x)}{SC(x')} &\leq \frac{\lambda}{1 - \mu}. \end{aligned}$$

Affine Congestion Games and Exact PoA

1. For congestion games with affine latencies:

$$\ell_r(|x|_r) = \alpha_r |x|_r + \beta_r.$$

2. Every such game is $(\frac{5}{3}, \frac{1}{3})$ -smooth, giving:

$$PoA \leq \frac{5}{2} = 2.5.$$

3. Proof (outline):

- Social cost:

$$SC(x) = \sum_i J_i(x) = \sum_r |x|_r \ell_r(|x|_r).$$

- Smoothness condition:

$$\sum_i J_i(x'_i, x_{-i}) \leq \frac{5}{3} SC(x') + \frac{1}{3} SC(x).$$

- Goal:

$$|x'_r| \ell_r(|x|_r + 1) \leq \frac{5}{3} |x'_r| \ell_r(|x'_r|) + \frac{1}{3} |x_r| \ell_r(|x|_r), \quad \forall x, x', r.$$

Congestion Pricing

- Design tolls $\tau_r(|x|_r)$ for each resource to incentivize optimal behavior:

$$\ell_r(|x|_r) + \tau_r(|x|_r).$$

12 Auctions

Utility of a Bidder

$$\text{Utility Bidder } i = \begin{cases} 0, & \text{if loses auction,} \\ v_i - p, & \text{if wins auction,} \end{cases}$$

where v_i is the bidder's valuation and p is the price paid.

Sealed-Bid Auction

- Each bidder privately submits a bid to the auctioneer.
- The auctioneer:
 1. Decides who gets the good (e.g., highest bid wins).
 2. Sets the selling price (affects bidder behavior).
- Bidders underbid because bidding true valuation yields zero utility. The degree of underbidding depends on unseen competitor bids.

Second-Price Auction (Vickrey)

- In a second-price auction, every bidder has a dominant strategy: bid their true valuation.
- Key properties:
 1. Dominant-Strategy Compatible (DSIC): Bidding truthfully is always the best strategy.
 2. Surplus Maximization: Bidding truthfully maximizes the social surplus:

$$\sum_{i=1}^n v_i x_i,$$

where $x_i = 1$ if bidder i wins.

3. Efficient Implementation: Can be implemented in linear time.
- Truthful bidding avoids the following risks:
 1. Overbidding: May result in winning at a price greater than valuation, causing negative utility.
 2. Underbidding: May result in losing an auction even when the item's value exceeds the second-highest bid.
 - Utility of a truth-telling bidder:

$$U_i = \begin{cases} v_i - B, & \text{if wins (where } B \text{ is the second-highest bid),} \\ 0, & \text{if loses.} \end{cases}$$

Sponsored Search Auctions

- Goals:
 1. DSIC: Truthful bidding should be dominant and yield non-negative utility.
 2. Surplus Maximization: Truthful bidding should maximize social surplus.
 3. Polynomial-Time: Assignments and payments should be computable efficiently.
- Process:
 1. Assign slots to bidders in descending order of bids to maximize surplus.
 2. Set prices to ensure DSIC behavior.

Myerson's Lemma

- Setting:

1. Allocation Space \mathcal{X} :
 - Single-item auctions: $\mathcal{X} = \{0, 1\}^n$, with $\sum_i x_i = 1$.
 - Sponsored search: $x_i = \alpha_j$ if bidder i is assigned slot j .
2. Allocation Rule $x(b) : \mathbb{R}_{\geq 0}^n \rightarrow \mathcal{X}$:
 - Single-item auctions: Allocate the good to the highest bid.
 - Sponsored search: Assign slot j to the j -th highest bid.
3. Payment Rule $p(b) : \mathbb{R}_{\geq 0}^n \rightarrow \mathbb{R}_{\geq 0}^n$:
 - Single-item auctions: Second-price payment.
 - Sponsored search: Determined via Myerson's Lemma.
4. Utility:

$$U_i(b) = v_i x_i(b) - p_i(b).$$

- Myerson's Lemma:

1. An allocation rule x is implementable if and only if it is monotone:

$$b'_i \geq b_i \implies x_i(b'_i, b_{-i}) \geq x_i(b_i, b_{-i}), \quad \forall b'_i, b_i, b_{-i}, \forall i.$$

2. For monotone x , there exists a unique payment rule p such that (x, p) is DSIC.
3. The payment rule can be derived analytically.

- Observations:

- Monotonicity ensures implementability.
- The payment rule is unique.

Proof of Monotonicity and Payment Rule

- DSIC implies that truthful reporting is the best strategy:

$$z \cdot x(z) - p(z) \geq z \cdot x(y) - p(y),$$

and

$$y \cdot x(y) - p(y) \geq y \cdot x(z) - p(z).$$

- Combining, the "sandwich inequality" holds:

$$z \cdot [x(y) - x(z)] \leq p(y) - p(z) \leq y \cdot [x(y) - x(z)].$$

- Implications:

- $x(z)$ must be monotone; otherwise, the inequality is violated.
- The payment difference $p(y) - p(z)$ is bounded by changes in allocation probability scaled by z and y .

Guessing the Price Function

- Use discontinuities in x to derive p :

$$z \cdot (x(y) - x(z)) \leq \Delta p \leq y \cdot (x(y) - x(z)).$$

- At a discontinuity:

$$\Delta p = h \cdot z,$$

where h is the height of the discontinuity.

- Fix $p(0) = 0$. Candidate formula:

$$p_i(b_i, b_{-i}) = \sum_{j=1}^l z_j \cdot (\text{jump in } x_i(\cdot, b_{-i}) \text{ at } z_j).$$